

# Regularized Logistic Regression sklearn

December 3, 2020

```
[100]: from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn import metrics
import pandas as pd
import numpy as np
```

```
[67]: data = pd.read_csv("./Data/caesarian.csv")
data.iloc[0:3,:]
```

```
[67]:
```

	Age	Delivery number	Delivery time	Blood of Pressure	Heart Problem	\
0	22	1	0	2	0	
1	26	2	0	1	0	
2	26	2	1	1	0	

  

	Caesarian
0	0
1	1
2	0

```
[125]: def test(df, split, penalty="l2"):
    data = pd.get_dummies(df, drop_first=True)
    X = data.iloc[:,0:len(data.columns)-1]
    y = data.iloc[:,len(data.columns)-1]
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
    random_state=42)
    lr = LogisticRegression(penalty=penalty,max_iter=10000).fit(X_train,
    y_train)
    accuracy = sum(y_test==lr.predict(X_test))/len(y_test)
    s = lr.score(X_test, y_test)
    auc = metrics.roc_auc_score(y_test, lr.predict_proba(X_test)[:,-1])
    return lr, auc, accuracy
```

```
[123]: data = pd.get_dummies(data, drop_first=True)
X = data.iloc[:,0:len(data.columns)-1]
y = data.iloc[:,len(data.columns)-1]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
random_state=42)
```

```
lr = LogisticRegression(penalty="l2",max_iter=10000).fit(X_train, y_train)
accuracy = sum(y_test==lr.predict(X_test))/len(y_test)
s = lr.score(X_test, y_test)
auc = metrics.roc_auc_score(y_test, lr.predict_proba(X_test)[:,-1])
```

```
[187]: data = pd.read_csv("./Data/caesarian.csv")
```

```
[189]: lg, auc, acc = test(data, 0.75, penalty = "none")
print("AUC:",auc,"Accuracy:",acc)
```

AUC: 0.7135416666666667 Accuracy: 0.7

```
[190]: lg, auc, acc = test(data, 0.75)
print("AUC:",auc,"Accuracy:",acc)
```

AUC: 0.6927083333333333 Accuracy: 0.7

```
[191]: data = pd.read_csv("Data/monks-1.test",delimiter=' ',header=None)
data = data.drop([0,8],axis=1)
cols = data.columns.tolist()
cols = cols[+1:] + cols[:+1]
data = data[cols]
```

```
[192]: lg, auc, acc = test(data, 0.75, penalty = "none")
print("AUC:",auc,"Accuracy:",acc)
```

AUC: 0.7347781217750259 Accuracy: 0.7129629629629629

```
[193]: lg, auc, acc = test(data, 0.75)
print("AUC:",auc,"Accuracy:",acc)
```

AUC: 0.7347781217750258 Accuracy: 0.7129629629629629

```
[194]: data = pd.read_csv("Data/credit-screening/crx.data")
```

```
[195]: lg, auc, acc = test(data, 0.75, penalty = "none")
print("AUC:",auc,"Accuracy:",acc)
```

AUC: 0.8029221653649339 Accuracy: 0.8034682080924855

```
[197]: lg, auc, acc = test(data, 0.75)
print("AUC:",auc,"Accuracy:",acc)
```

AUC: 0.920549420953407 Accuracy: 0.8439306358381503

```
[198]: data = pd.read_csv("Data/framingham.csv")
```

```
[199]: lg, auc, acc = test(data, 0.75, penalty = "none")
print("AUC:",auc,"Accuracy:",acc)
```

AUC: 0.9357530864197531 Accuracy: 0.8765027322404372

```
[200]: lg, auc, acc = test(data, 0.75)
print("AUC:",auc,"Accuracy:",acc)
```

AUC: 0.9360493827160494 Accuracy: 0.8797814207650273

```
[204]: data = pd.read_csv("Data/default of credit card clients.csv")
```

```
[205]: pd.get_dummies(data, drop_first=True)
```

```
[205]:
```

	ID	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_0	PAY_2	PAY_3	\
0	1	20000	2	2	1	24	2	2	-1	
1	2	120000	2	2	2	26	-1	2	0	
2	3	90000	2	2	2	34	0	0	0	
3	4	50000	2	2	1	37	0	0	0	
4	5	50000	1	2	1	57	-1	0	-1	
...	...	...	...	...	...	...	...	...	...	
29995	29996	220000	1	3	1	39	0	0	0	
29996	29997	150000	1	3	2	43	-1	-1	-1	
29997	29998	30000	1	2	2	37	4	3	2	
29998	29999	80000	1	3	1	41	1	-1	0	
29999	30000	50000	1	2	1	46	0	0	0	

	PAY_4	...	BILL_AMT4	BILL_AMT5	BILL_AMT6	PAY_AMT1	PAY_AMT2	\
0	-1	...	0	0	0	0	689	
1	0	...	3272	3455	3261	0	1000	
2	0	...	14331	14948	15549	1518	1500	
3	0	...	28314	28959	29547	2000	2019	
4	0	...	20940	19146	19131	2000	36681	
...	...	...	...	...	...	...	...	
29995	0	...	88004	31237	15980	8500	20000	
29996	-1	...	8979	5190	0	1837	3526	
29997	-1	...	20878	20582	19357	0	0	
29998	0	...	52774	11855	48944	85900	3409	
29999	0	...	36535	32428	15313	2078	1800	

	PAY_AMT3	PAY_AMT4	PAY_AMT5	PAY_AMT6	default payment next month
0	0	0	0	0	1
1	1000	1000	0	2000	1
2	1000	1000	1000	5000	0
3	1200	1100	1069	1000	0
4	10000	9000	689	679	0
...	...	...	...	...	...

29995	5003	3047	5000	1000	0
29996	8998	129	0	0	0
29997	22000	4200	2000	3100	1
29998	1178	1926	52964	1804	1
29999	1430	1000	1000	1000	1

[30000 rows x 25 columns]

```
[202]: lg, auc, acc = test(data, 0.75, penalty = "none")
print("AUC:",auc,"Accuracy:",acc)
```

AUC: 0.6501707782984041 Accuracy: 0.7830666666666667

```
[203]: lg, auc, acc = test(data, 0.75)
print("AUC:",auc,"Accuracy:",acc)
```

AUC: 0.6501443010428375 Accuracy: 0.7830666666666667

```
[ ]:
```